

Projective Reed–Solomon Syndromes Beyond Redundancy Four: Complete First Levels and Coherent Polar Flags

Tavis Rudd

Abstract

We classify, up to projective-semilinear equivalence, the one-column maximum-distance-separable extensions of projective Reed–Solomon codes at redundancy five over every prime power $q \geq 7$, and give exact projective counts. Equivalently, we classify the deep-hole syndrome directions in the first case left open by the classification through redundancy four. A syndrome determines a Hankel linear system of binary forms and is deep exactly when that system contains no completely split squarefree member. The generic redundancy-five case is controlled by an integral off-diagonal fiber-square curve of arithmetic genus one.

For higher redundancy we introduce coherent polar flags: iterated contractions in which every selected factor remains a marked forbidden root. Their basic identity is $g \in W_{\iota_{\lambda}f} \Rightarrow \lambda g \in W_f$, with a squarefree lift exactly when $\lambda \nmid g$. Separating contained components from transverse point counts gives complete all-field syndrome classifications at redundancies six and seven; the redundancy-seven split-free list is a deep-hole classification whenever the covering radius is six, in particular for $q \geq 11$. More generally, once the contained-component assertions and the explicit modular pullback-degree bound are known through redundancy r , the transverse branch has the linear-scale threshold $q \geq 6r - 15 + \lfloor 2\sqrt{6r - 17} \rfloor$.

Index Terms

Algebraic coding theory, deep holes, finite fields, maximum-distance-separable extensions, normal rational curves, projective Reed–Solomon codes.

I. Introduction

This paper has two principal contributions. First, it gives the complete all-field classification at redundancy five, the first case left open by the known classification through redundancy four. Second, it introduces coherent polar flags: a marked contraction method that propagates splitting information to higher redundancies without forgetting the factor needed for a squarefree lift. This gives complete all-field syndrome classifications at redundancies six and seven. We use deep-hole syndromes because their Hankel systems make both the base classification and the propagation mechanism visible.

The mechanism is the elementary identity

$$g \in W_{\iota_{\lambda}f} \iff \lambda g \in W_f, \quad \lambda g \text{ squarefree} \iff \lambda \nmid g$$

when g is squarefree. Retaining λ is therefore the essential coherence condition. Conditional on the contained-component assertions and the modular pullback-degree bound in Theorem V.10, the transverse branch closes at redundancy r for

$$q \geq 6r - 15 + \lfloor 2\sqrt{6r - 17} \rfloor.$$

Let $\text{PRS}(k)$ denote the projective Reed–Solomon code of length $q + 1$ obtained by evaluating homogeneous binary forms of degree $k - 1$ on $\mathbb{P}^1(\mathbb{F}_q)$. If the redundancy is

$$r = q + 1 - k,$$

a parity-check matrix has as columns the $\mathbb{F}_q$ -points of the degree- $(r - 1)$ normal rational curve

$$\mathcal{R}_{r-1} = \{(1, t, \dots, t^{r-1}) : t \in \mathbb{F}_q\} \cup \{(0, \dots, 0, 1)\} \subset \mathbb{P}^{r-1}.$$

Call a projective syndrome direction split-free if it is not contained in the span of $r - 2$ distinct columns. When $\rho(\text{PRS}(q + 1 - r)) = r - 1$, the split-free directions are precisely the projective syndrome directions of deep-hole cosets. The logical separation used throughout the paper is

$$\boxed{\text{split-free classification} + \rho(\text{PRS}(q + 1 - r)) = r - 1 \implies \text{deep-hole classification.}}$$

The first term is proved geometrically; the covering-radius equality is a separate coding-theoretic input and is never inferred from a finite split-free census.

The correspondence between deep holes, one-symbol MDS extensions, and uncovered points of a normal rational curve was developed in [1], [2]. Zhang, Wan, and Kaipa classified projective Reed–Solomon deep-hole directions through

redundancy four and identified redundancy five as the next case [2]. Our first theorem gives the corresponding all-field redundancy-five classification.

Theorem I.1 (Complete redundancy-five classification). Let $q \geq 7$ be a prime power. The covering radius of $\text{PRS}(q-4)$ is 4. Its projective deep-hole syndromes are exactly the following disjoint families:

- (i) the tangent family T , of size $q(q+1)$;
- (ii) the conjugate-secant family S , of size $q(q^2-1)/2$;
- (iii) if $\text{char } \mathbb{F}_q \neq 3$, the rational osculating-pair family O^+ when $q \equiv 2 \pmod{3}$, of size $q(q+1)/2$, or the conjugate osculating-pair family O^- when $q \equiv 1 \pmod{3}$, of size $q(q-1)/2$;
- (iv) in characteristic three, one $\text{PGL}_2(q)$ -fixed nucleus point and one wild Artin-Schreier orbit W of size $(q^2-1)/2$;
- (v) the sporadic $\text{PGL}_2(q)$ -orbits in Table IV, occurring exactly for $q \in \{7, 8, 9, 11, 13, 17, 19\}$.

There are no other deep-hole syndromes. The three corresponding nonsporadic totals are

$$\frac{q^2(q+3)}{2}, \quad \frac{q(q+1)(q+2)}{2}, \quad \frac{q^3+3q^2+q+1}{2},$$

according as $q \equiv 1 \pmod{3}$ with characteristic not three, $q \equiv 2 \pmod{3}$, or $\text{char } \mathbb{F}_q = 3$.

The main mechanism is a characteristic-free Hankel dictionary. A syndrome at redundancy r annihilates a linear system of binary forms of degree $r-2$. It is split-free if and only if this system contains no completely split squarefree form; it is a deep-hole direction when, in addition, $\rho(\text{PRS}(q+1-r)) = r-1$. At redundancy five the system is a pencil of cubics. Its induced degree-three map of projective lines is either cyclic or has geometric monodromy S_3 . Cyclic descent gives the arithmetic families in Theorem I.1; in the S_3 case the off-diagonal fiber square is an integral curve of arithmetic genus one, and the singular-curve Weil bound [3] forces a split fiber for $q \geq 23$. The certified finite computations in Section VII handle the remaining fields.

Starting at redundancy six, the Hankel system is no longer a pencil. The useful replacement is the coherently parameterized line of contractions obtained by choosing one prospective root. Iteration produces a polar flag. Geometrically, naïve contraction keeps the lower divisor but forgets which point of $\mathbb{P}^1$ was removed. If the lower divisor passes through that point, putting the point back creates a double point rather than a squarefree divisor. The polar flag retains the removed point as a forbidden marker, exactly the datum needed for lifting. This pointed coherence is essential: at $q = 19$ an individual contracted fiber has no split witness avoiding its marker, but no consecutive-row polar line is contained in its orbit.

The resulting structural theorem has two outputs.

Theorem I.2 (Classification results and explicit gates beyond five). The following statements hold.

- (a) For every prime power $q \geq 7$, the projective deep-hole directions of $\text{PRS}(q-5)$ are exactly the persistent tangent and conjugate-secant families, the exceptional orbits listed for $q = 7, 8, 9, 11, 13$, and, when $q = 2^m$ with odd $m \geq 5$, one characteristic-two nucleus orbit.
- (b) The displayed finite domains through $q = 32$ have the certified split-free classification stated in Section VI. Together with the structural theorem, these finite classifications give the all-field result: for every $q \geq 13$ the split-free locus consists of the persistent families and, when $q = 2^m$ with m odd, one central nucleus point. This is a deep-hole classification only when $\rho(\text{PRS}(q-6)) = 6$, in particular for $q \geq 11$.

At redundancy r , the semilinear orbits in the persistent conjugate-secant family are parameterized by

$$T/T^{r-1} \quad \text{modulo inversion and Frobenius,}$$

where $T = \{z \in \mathbb{F}_{q^2}^\times : z^{q+1} = 1\}$ is the norm-one torus and Frobenius is coefficient Frobenius. On the tangent family the unipotent coordinate transforms as $z \mapsto z + (r-1)u$; consequently its orbit decomposition changes precisely when $\text{char } \mathbb{F}_q$ divides $r-1$.

The two parts have different scopes. Both are unconditional all-field syndrome classifications, but part (b) becomes a deep-hole classification only when $\rho(\text{PRS}(q-6)) = 6$. Table II places the split-free result, radius gate, deep-hole conclusion, and remaining field gap in separate columns for each redundancy.

a) Relation to previous work.: The tangent and conjugate-secant families and the coding/normal-rational curve dictionary come from [1], [2]; the covering-radius input in the high-rate range is due to Seroussi and Roth [4]. The broader arc/MDS and code-automorphism background is surveyed or established in [5], [6]. Earlier explicit-family results for projective and generalized projective Reed-Solomon codes include [7], [8], [9]; the recent extension framework of [10] further clarifies the deep-hole/MDS-extension interface. These works construct families, treat punctured or lower-redundancy regimes, or analyze the extension criterion; they do not supply the complete redundancy-five cubic-pencil exhaustion proved here. Classical apolarity, catalecticants, and binary-form rank strata are treated in [11], [12]; our use of divided powers keeps the small-characteristic contraction formulas explicit. From the viewpoint of binary forms,

the question is which linear systems contain a totally split squarefree member. Factorization statistics in linear families [13] and normal-rational-curve nuclei [14] are inputs. For splitting families, we use Wang’s Theorem 3.6 [15]: on the étale locus, the factorization type of a specialized polynomial is the cycle type of its Frobenius conjugacy class. The coding-theory references above identify deep syndromes with uncovered points of a normal rational curve and settle the lower-redundancy families; they do not classify the redundancy-five cubic pencils. The orbit-theoretic and factorization references supply normal forms and distribution tools, but not the coherent marked contraction used here. The nucleus results locate characteristic-dependent kernels, while the present work computes which of their coherent lifts meet the PRS splitting incidence. Accordingly, this paper’s contributions are separated into the redundancy-five through seven classifications and the conditional transverse induction mechanism. Claim-by-claim novelty and dependency boundaries are recorded in the electronic ledger.

b) Organization.: Section II first gives a proof-free guide to the argument, and Section III gives the Hankel dictionary. Section IV proves Theorem I.1. Sections V and VI develop coherent polar induction and prove Theorem I.2. Section VII states the computational and formal trust boundary, and Section VIII records the exact open boundary.

II. Guide to the proof

This section is a proof-free guided tour. Its purpose is to install the information flow used throughout the paper before coordinates, finite exceptional tables, or formal trust boundaries intervene.

a) Syndrome to splitting system.: A syndrome direction $[f] \in \mathbb{P}(\Gamma^{r-1}E)$ determines its Hankel annihilator $W_f \subset \text{Sym}^{r-2}E^\vee$. A product of $r - 2$ distinct rational linear factors lies in W_f exactly when f lies in the span of the corresponding normal-rational-curve columns. Thus the geometric classification begins with split-free systems. The separate covering-radius theorem is what promotes split-free directions to code deep holes; the dictionary alone does not.

b) The redundancy-five base cover.: At redundancy five, W_f is a pencil of binary cubics. Its root incidence is a degree-three map of projective lines. The off-diagonal fiber square separates the cyclic cases, controlled by Frobenius on a deck transformation, from the S_3 case, controlled by rational points on a geometrically integral $(2, 2)$ curve. This is the bottom cover used at every later level.

c) Why independent lower fibers fail.: At higher redundancy, choosing one lower witness forgets which factor was removed. If that lower witness contains the chosen factor, then multiplication restores it with multiplicity two rather than producing a squarefree form. The lost datum is therefore not auxiliary: it is the exact obstruction to lifting.

d) Coherent replacement and the fork.: Choose a prospective root λ , contract f to $\iota_\lambda f$, and retain λ as a marked forbidden root. Iteration gives an ordered polar flag. Its first-polar line has two qualitatively different behaviors. If it is transverse to the lower bad locus and collision schemes, only finitely many parameters must be removed and a point count produces an escape. If it is scheme-theoretically contained, the point-count argument says nothing: ordinary catalecticant components and components created in small characteristic require a separate level-specific classification. Section V introduces the persistent and modular loci only when that distinction enters the proof.

e) Lifting.: On the transverse branch, a rational point of the lower ordered splitting cover outside the branch, diagonal, marker, and collision divisors produces a split squarefree lower witness avoiding all markers. Multiplying by the marked factors then lifts it squarefreely to W_f . Figure 1 records this sequence at the point where the construction is formalized.

TABLE I
Mathematical dependency map.

Result	Printed geometric argument	Imported input
R5	Hankel dictionary; gcd strata; cyclic descent; integral $(2, 2)$ fiber square and point bound	lower tangent/secant families; high-rate radius theorem; singular-curve bound
R6	one-marker polar line; ordinary and binary contained components; transverse escape	R5 lower package and radius gate
R7	two-marker contraction; complete second-marker bad scheme, including cyclic and fixed-gcd collisions; contained rank-two theorem; transverse escape	R5/R6 lower packages; separate covering-radius result

TABLE II

Classification status by redundancy. “Unconditional” refers to the displayed field range. A split-free classification becomes a deep-hole classification only after the separate radius gate.

Red.	Field range	Split-free geometry	Deep holes/MDS extensions	Remaining gap
5	$q \geq 7$	Complete, unconditional	Complete, unconditional; $\rho = 4$	None
6	$q \geq 7$	Complete, unconditional	Complete, unconditional; $\rho = 5$	None
7	$q \geq 7$	Complete, unconditional	Complete for $q \geq 11$; conditional on $\rho = 6$ at $q = 7, 8, 9$	Radius at $q = 7, 8, 9$

III. The syndrome–Hankel dictionary

Let E be a two-dimensional vector space. We use divided powers for syndromes and ordinary symmetric powers for root forms. This avoids invalid binomial rescalings in small characteristic. For $f = (a_0 : \cdots : a_{r-1}) \in \mathbb{P}(\Gamma^{r-1}E)$, write $e_0, \dots, e_{r-1}$ for the divided-power coordinate basis and $\nu(t) = (1, t, \dots, t^{r-1})$, with $\nu(\infty) = e_{r-1}$. If x, y is a basis of E with dual basis X, Y , our perfect divided-power/symmetric-power pairing is

$$\left\langle x^{[n-i]} y^{[i]}, X^{n-j} Y^j \right\rangle = \delta_{ij}, \quad \left\langle \sum_{i=0}^n a_i x^{[n-i]} y^{[i]}, \sum_{i=0}^n b_i X^{n-i} Y^i \right\rangle = \sum_{i=0}^n a_i b_i.$$

Thus the pairing is coefficient extraction, with no binomial factors in any characteristic. Define

$$H_f = \begin{pmatrix} a_0 & a_1 & \cdots & a_{r-2} \\ a_1 & a_2 & \cdots & a_{r-1} \end{pmatrix}, \quad W_f = \ker H_f \subset \text{Sym}^{r-2} E^\vee.$$

TABLE III

Notation used from the Hankel dictionary onward.

Symbol	Meaning
q	field size; all numerical thresholds are thresholds on prime powers
k	code dimension in $\text{PRS}(k)$
$r = q + 1 - k$	redundancy
$n = r - 1$	divided-power degree of a syndrome representative
E	a two-dimensional vector space over the current base field
$f \in \Gamma^n E$	divided-power representative of a projective syndrome
$W_f \subset \text{Sym}^{n-1} E^\vee$	Hankel/apolar witness system annihilated by f
g, δ, κ	genus, deletion degree, and point-bound correction
d	degree of the unavailable marker scheme

Read $f = (a_0, \dots, a_{r-1})$ as a finite sequence. A form $g \in W_f$ is the characteristic polynomial of a recurrence satisfied by that sequence, while a squarefree split g has a basis of geometric-sequence solutions. Thus split-free means that the syndrome sequence satisfies no order- $(r-2)$ recurrence with squarefree completely split characteristic polynomial.

Lemma III.1 (Hankel criterion). Let $\lambda_1, \dots, \lambda_{r-2}$ be distinct points of $\mathbb{P}(E^\vee)(\mathbb{F}_q)$ and put $g = \prod_i \lambda_i$. Then

$$f \in \text{span}(\nu(\lambda_1), \dots, \nu(\lambda_{r-2})) \iff g \in W_f.$$

The same statement with repeated factors and Hasse derivatives gives the corresponding osculating spans. Both statements commute with $\text{PGL}_2(q)$.

Proof. On an affine chart, the geometric sequences $\nu(\lambda_i)$ solve the recurrence with characteristic polynomial g ; their independence is the Vandermonde determinant. A root at infinity gives the reversed-chart recurrence. For repeated roots, the confluent Vandermonde basis is formed by Hasse derivatives, so the same argument works in every characteristic. Equivariance follows because substitution transports both the normal rational curve and the recurrence. $\square$

Corollary III.2 (Split-free criterion). Assume $\rho(\text{PRS}(q+1-r)) = r-1$. A projective syndrome f is deep if and only if W_f contains no completely split squarefree form of degree $r-2$ over $\mathbb{F}_q$.

We use split-free without a covering-radius hypothesis and call a direction shallow when W_f contains a completely split squarefree form.

The first locus that persists under contraction is already visible here. If every member of W_f has a common quadratic factor, then that factor is either a rational square, an irreducible quadratic, or a product of two distinct rational factors.

These are respectively the tangent deep family, the conjugate-secant deep family, and a shallow rational secant. The common factor cannot grow with redundancy: since $\dim W_f = r - 3$,

$$r - 3 \leq r - 1 - \deg \gcd(W_f),$$

so its degree is at most two.

A. Coding consequences and recognition

The dictionary has two immediate coding consequences. First, a split-free direction is an uncovered point of the normal rational curve and hence gives a one-column MDS extension whenever the stated covering-radius gate applies. The classification theorems therefore give exact projective deep-hole counts, not only normal forms. Second, they give recognition procedures. From a syndrome one forms H_f , computes W_f , and tests the listed geometric families: a common quadratic factor detects the persistent tangent and conjugate-secant families; the relevant contraction kernel detects the modular locus; and the finite exceptional fields are decided by the canonical orbit representatives printed in Table IV and recorded in the supplement.

Recognition at the fixed redundancies proved here begins with row-reduction of H_f and the listed family-membership tests. Coherent polar induction replaces a direct search through the projective members of W_f .

IV. Redundancy five: exceptional cubic covers

Here $f = (a_0 : \dots : a_4)$ and W_f is the pencil annihilated by

$$\begin{pmatrix} a_0 & a_1 & a_2 & a_3 \\ a_1 & a_2 & a_3 & a_4 \end{pmatrix}.$$

These are divided-power coordinates; the coefficients of a cubic in W_f are ordinary symmetric-power coordinates. Factoring the formal quartic $\sum a_i T^i U^{4-i}$ would therefore be noninvariant in characteristics two and three.

Proposition IV.1 (Radius gate). For every prime power $q \geq 7$,

$$\rho(\text{PRS}(q-4)) = 4.$$

Proof. The coding/arc criterion identifies radius four with completeness of the quartic normal rational curve. Seroussi–Roth prove completeness when

$$k \geq \left\lfloor \frac{q-1}{2} \right\rfloor$$

apart from the even-field dimensions $k = 2, q-2$ [4]. Here $k = q-4$, the inequality holds for $q \geq 7$, and neither exceptional dimension occurs. Thus Corollary III.2 applies. The finite certificate independently checks the same conclusion only in its declared finite domain; it is not the all-field radius proof. $\square$

A. Gcd strata

Proposition IV.2 (Quadratic gcd). If $\deg \gcd(W_f) = 2$, then

$$W_f = Q\langle T, U \rangle.$$

If Q is split squarefree, f is shallow. If Q is irreducible, f belongs to the conjugate-secant family; if $Q = \lambda^2$, it belongs to the tangent family. Conversely every point in those families has the displayed gcd. Their deep-family sizes are

$$|T| = q(q+1), \quad |S| = \frac{q(q^2-1)}{2}.$$

Proof. The equality follows by dimension. A split Q and a linear cofactor avoiding its two roots give a split squarefree cubic. An irreducible Q survives in every member, while a square Q forces a repeated root in every member. Lemma III.1, including its confluent and conjugate-root forms, identifies the corresponding syndromes with the secant and tangent spans. The sizes are the established tangent/conjugate-secant counts from the lower-redundancy classification [1], [2]. $\square$

Proposition IV.3 (Linear gcd). Suppose $\deg \gcd(W_f) = 1$. Then f is shallow for $q \geq 8$. At $q = 7$ the same conclusion is part of Certificate R5.

Proof. Write $W_f = \lambda_0 V$ with V a basepoint-free quadratic pencil. If its degree-two map $\psi : \mathbb{P}^1 \rightarrow \mathbb{P}^1$ is separable, its rational deck involution ι gives the graph

$$\Gamma_\iota(\mathbb{F}_q) = \{(t, \iota(t)) : t \in \mathbb{P}^1(\mathbb{F}_q)\},$$

which has exactly $q + 1$ rational points. A safe union bound deletes at most two fixed graph points, at most four graph points over the branch fibers, and at most the two graph points $(\lambda_0, \iota(\lambda_0))$ and $(\iota(\lambda_0), \lambda_0)$. Thus fewer than $q + 1$ graph points are deleted when $q \geq 8$. A surviving point gives a split quadratic with two distinct roots avoiding λ_0 , and hence a split squarefree cubic in W_f . In characteristic two an inseparable quadratic pencil is equivalent to $\langle T^2, U^2 \rangle$; the two overlapping Hankel equations then force f onto the normal rational curve, contrary to the rank-two pencil hypothesis. $\square$

B. The trivial-gcd cubic cover

Choose a basis c_1, c_2 of a trivial-gcd pencil and put

$$X_f = \{([u : v], t) : uc_1(t) + vc_2(t) = 0\} \subset \mathbb{P}^1 \times \mathbb{P}^1.$$

Proposition IV.4 (Cubic incidence and off-diagonal fiber square). The curve X_f is geometrically integral of bidegree $(1, 3)$ and arithmetic genus zero. Projection to the root coordinate is birational, while projection to the pencil coordinate gives a degree-three morphism

$$\phi_f : \mathbb{P}^1 \longrightarrow \mathbb{P}^1.$$

If ϕ_f is separable, division of

$$c_1(t_1)c_2(t_2) - c_2(t_1)c_1(t_2)$$

by the diagonal bracket defines a residual curve $Y_f \subset \mathbb{P}^1 \times \mathbb{P}^1$ of bidegree $(2, 2)$ and arithmetic genus one. Its geometric components are the geometric-monodromy orbits on ordered pairs of distinct roots.

Proof. A component of X_f of bidegree $(0, e)$ would be a common factor of c_1, c_2 ; hence none exists. The bidegree genus formula gives genus zero. A root determines the unique pencil member containing it, so the root projection is birational. For a separable map the diagonal factor in the fiber-square determinant has multiplicity one. Removing its bidegree $(1, 1)$ from the fiber square leaves bidegree $(2, 2)$, whose arithmetic genus is one. The standard fiber-square/monodromy correspondence identifies its components with orbits on distinct ordered sheets. $\square$

Lemma IV.5 (Cyclic stratum). Suppose ϕ_f is geometrically cyclic.

- (i) In characteristic different from three there are two ramification points. If they are rational, the corresponding osculating-plane intersection is deep exactly when $q \equiv 2 \pmod{3}$; if they are conjugate, it is deep exactly when $q \equiv 1 \pmod{3}$.
- (ii) In characteristic three all osculating planes meet in $n = (0 : 0 : 1 : 0 : 0)$, whose pencil is $\langle T^3, U^3 \rangle$; this fixed point is deep.
- (iii) The remaining wild characteristic-three pencils have the form $\langle U^3, T^3 + aTU^2 \rangle$. They are deep exactly when $-a$ is nonsquare, forming one orbit of size $(q^2 - 1)/2$.

Proof. Assume first that the characteristic is not three. Riemann–Hurwitz gives two totally ramified points t_1, t_2 . The confluent Hankel criterion places f in $\text{Osc}(t_1) \cap \text{Osc}(t_2)$; after transporting the pair to $\{0, \infty\}$, the pencil is $\langle T^3, U^3 \rangle$ and a geometric deck generator is $t \mapsto \zeta t$.

For a rational ramification pair, Frobenius fixes the points and the deck map descends exactly when $\zeta^q = \zeta$, namely $q \equiv 1 \pmod{3}$. Thus its complement is deep exactly for $q \equiv 2 \pmod{3}$. The setwise stabilizer of the ordered normal form is the split dihedral group of order $2(q - 1)$, so the orbit size is

$$\frac{q(q^2 - 1)}{2(q - 1)} = \frac{q(q + 1)}{2}.$$

For a conjugate pair Frobenius interchanges the fixed points and inverts the multiplier; descent is instead $q \equiv 2 \pmod{3}$. The nonsplit dihedral stabilizer has order $2(q + 1)$ and the orbit size is $q(q - 1)/2$.

Now suppose the characteristic is three. The second Hasse derivative of

$$\nu(t) = (1, t, t^2, t^3, t^4)$$

is $(0, 0, 1, 3t, 6t^2) = e_2$, including in the reversed infinity chart. Thus every osculating plane contains $n = [e_2]$. Substitution shows that n is $\text{PGL}_2(q)$ -fixed and $W_n = \langle T^3, U^3 \rangle$, so it is deep.

A remaining separable cyclic cover has one wild totally ramified point. Sending it to infinity and comparing coefficients under a nontrivial translation gives

$$W_f = \langle U^3, T^3 + \alpha T^2 U + \beta T U^2 \rangle, \quad \alpha = 0, \quad \kappa^2 = -\beta.$$

It is therefore represented by

$$f_a = e_2 - ae_4, \quad W_{f_a} = \langle U^3, T^3 + aTU^2 \rangle.$$

The affine members are $z^3 + az + s$. Such a member has three rational roots exactly when the additive map $z \mapsto z^3 + az$ has a nonzero rational kernel, equivalently when $-a$ is a square. Hence the nonsquare parameters form the deep orbit. Its stabilizer is $t \mapsto \pm t + \beta$, of order $2q$, and orbit-stabilizer gives $(q^2 - 1)/2$. $\square$

Lemma IV.6 (S_3 stratum). If ϕ_f has geometric monodromy S_3 and $q \geq 23$, then the pencil contains a completely split squarefree cubic.

Proof. By Proposition IV.4, components of Y_f correspond to monodromy orbits on distinct ordered roots. The S_3 action is two-transitive, so Y_f is geometrically integral and has arithmetic genus one. The Aubry–Perret singular-curve bound, in the form

$$|\#C(\mathbb{F}_q) - (q + 1)| \leq 2\pi_C \sqrt{q}$$

for a reduced geometrically integral projective curve of arithmetic genus π_C , gives

$$\#Y_f(\mathbb{F}_q) \geq q + 1 - 2\sqrt{q}$$

[3, p. 468]. An integral curve of arithmetic genus one has at most one geometric singular point: if it is singular, its normalization has genus zero and the total delta invariant is one. We therefore delete at most one rational singular point. The diagonal costs

$$Y_f \cdot \Delta = (2, 2) \cdot (1, 1) = 4.$$

Riemann–Hurwitz gives different degree four for the separable cubic map, hence at most four branch values. A singular cubic has at most two distinct roots and at most two ordered off-diagonal pairs, so the branch deletion costs at most eight. If $q + 1 - 2\sqrt{q} > 13$, a smooth rational point remains outside the diagonal and branch fibers. It gives two distinct rational roots of a squarefree cubic, whose third root is rational. This inequality is equivalent to $q - 2\sqrt{q} > 12$ and holds for every prime power $q \geq 23$. $\square$

Proposition IV.7 (Exhaustion and finite bridge). The preceding cases exhaust every redundancy-five pencil for $q \geq 23$. Certificate R5 closes exactly

$$q \in \{7, 8, 9, 11, 13, 16, 17, 19\}.$$

It finds sporadic S_3 -stratum orbits exactly at $q = 7, 8, 9, 11, 13, 17, 19$ and none at $q = 16$.

Proof. The gcd degree is zero, one, or two. Propositions IV.2 and IV.3 settle the positive-gcd cases. In the trivial-gcd case the cubic map is inseparable or separable. The inseparable case occurs only in characteristic three and the overlap equations give the all-cube nucleus pencil. A separable transitive cubic has geometric monodromy C_3 or S_3 , settled by Lemmas IV.5 and IV.6. The stated finite residue, canonical representatives, stabilizers, histograms, and Frobenius links are exactly the domain of Certificate R5. $\square$

TABLE IV: Complete sporadic orbit inventory at redundancy five. A representative $[a_0, \dots, a_4]$ is in the divided-power syndrome coordinates of Section III. The last column records the nontrivial Frobenius fusion; a dash means the orbit is fixed.

q	representative	orbit size	stabilizer	Frobenius
7	[0,1,0,0,3]	28	12	–
7	[0,1,0,0,2]	112	3	–
7	[0,0,1,0,3]	168	2	–
7	[0,1,0,3,2]	168	2	–
7	[1,0,0,2,1]	168	2	–
8	[0,1,1,0,2]	252	2	$\rightarrow [0, 1, 1, 0, 4]$
8	[0,1,1,0,4]	252	2	$\rightarrow [0, 1, 1, 0, 6]$
8	[0,1,1,0,6]	252	2	$\rightarrow [0, 1, 1, 0, 2]$
9	[1,0,0,1,2]	180	4	–
9	[0,1,0,4,1]	360	2	$\leftrightarrow [0, 1, 0, 4, 4]$
9	[0,1,0,4,4]	360	2	$\leftrightarrow [0, 1, 0, 4, 1]$
11	[1,0,0,1,10]	330	4	–
11	[1,0,0,1,1]	660	2	–
13	[0,1,0,0,4]	182	12	–
13	[1,0,0,4,7]	546	4	–
17	[1,0,0,1,5]	1224	4	–

q	representative	orbit size	stabilizer	Frobenius
19	[0,1,0,0,4]	570	12	–

The certificate additionally records the complete cubic-pencil member histogram of each representative. Repeated orbit sizes in Table IV are therefore visibly distinguished without requiring a working-directory file.

TABLE V

Independent R5 completeness totals. The first equality is the direct syndrome enumeration versus the orbit–stabilizer sum; the second decomposes the same total into the formula in Theorem I.1 and the sporadic rows above.

q	direct = orbit sum	nonsporadic	sporadic
7	889 = 889	245	644
8	1116 = 1116	360	756
9	1391 = 1391	491	900
11	1848 = 1848	858	990
13	2080 = 2080	1352	728
16	2432 = 2432	2432	0
17	4131 = 4131	2907	1224
19	4541 = 4541	3971	570

V. Coherent polar induction

A. Spaces, markers, and bad schemes

Fix a field k , a two-dimensional k -space E , and

$$S_n = \mathbb{P}(\Gamma^n E).$$

For $\lambda \in E^\vee$, define divided-power contraction by

$$\langle \iota_\lambda f, g \rangle = \langle f, \lambda g \rangle.$$

Then

$$g \in W_{\iota_\lambda f} \iff \lambda g \in W_f. \quad (1)$$

The lower witness lifts squarefreely only when $\lambda \nmid g$. Thus the factor selected for contraction must remain as a marked forbidden root.

Definition V.1 (Persistent and modular loci). The persistent locus is the ordinary rank-two catalecticant locus propagated by contraction. The modular locus is the additional contraction-kernel locus that can appear in positive characteristic. A collision divisor records parameters at which the marked factor becomes a repeated root.

Definition V.2 (Contraction and marker spaces). Let $C_f : E^\vee \rightarrow \Gamma^{n-1} E$ be $\lambda \mapsto \iota_\lambda f$, and put $U_f = \mathbb{P}(E^\vee) \setminus \mathbb{P}(\ker C_f)$. The first-polar rational map is

$$\Phi_f : U_f \longrightarrow S_{n-1}, \quad [\lambda] \longmapsto [\iota_\lambda f].$$

Its image is a point or a line. For $j \geq 1$, let

$$\text{Conf}_j(\mathbb{P}(E^\vee)) = \{(\lambda_1, \dots, \lambda_j) : \lambda_i \neq \lambda_\ell \text{ for } i \neq \ell\}.$$

The ordered marker space $U_{f,j}$ is the open subset on which every successive contraction is nonzero. The universal j -step polar flag is the map

$$(\lambda_1, \dots, \lambda_j) \longmapsto ([\iota_{\lambda_1} f], \dots, [\iota_{\lambda_j} \cdots \iota_{\lambda_1} f]).$$

In an affine coordinate r the first map is

$$\Phi_f(r) = (a_1 - ra_0, \dots, a_n - ra_{n-1}), \quad \Phi_f(\infty) = (a_0, \dots, a_{n-1}).$$

Example V.3 (The complete R5-to-R6 propagation at $q = 29$). Work over $\mathbb{F}_{29}$ at redundancy six and take

$$f = (6 : 1 : 0 : 0 : 5 : 1).$$

The Hankel kernel is the net of quartics

$$W_f = \ker \begin{pmatrix} 6 & 1 & 0 & 0 & 5 \\ 1 & 0 & 0 & 5 & 1 \end{pmatrix} = \langle (0, 0, 1, 0, 0), (24, 1, 0, 1, 0), (28, 1, 0, 0, 1) \rangle.$$

The displayed basis has no common factor, so the persistent test does not decide the syndrome. Choose the first marker $r = 0$. Then

$$\iota_0 f = (1 : 0 : 0 : 5 : 1).$$

This is now a redundancy-five syndrome, and its full cubic pencil is

$$W_{\iota_0 f} = \langle (0, 1, 0, 0), (25, 0, 1, 24) \rangle.$$

The member with pencil parameter $[9 : 28]$ is

$$h(t) = 4 + 9t - t^2 + 5t^3 \in W_{\iota_0 f},$$

because the two Hankel products are $4 + 5 \cdot 5 = 0$ and $5(-1) + 5 = 0$ in $\mathbb{F}_{29}$. Direct factorization gives

$$h(t) = 5(t - 2)(t - 6)(t - 27).$$

None of these roots is the marked root 0. Equation (1) therefore lifts h to

$$t h(t) = (0, 4, 9, 28, 5) \in W_f,$$

a squarefree quartic with roots 0, 2, 6, 27. The Hankel criterion now places f in the span of the corresponding four normal-rational-curve columns, so f is shallow. This one calculation exhibits every interface in the R5-to-R6 propagation: the upper Hankel net, the ordinary common-factor test, contraction to the complete R5 cubic pencil, selection on its splitting cover, marker avoidance, and squarefree lifting. In the general proof the R5 curve argument supplies such a pencil parameter after deleting the marked fibers; the displayed parameter $[9 : 28]$ is one complete instance of that step.

Definition V.4 (One-step lower package). Fix $f \in S_n(\mathbb{F}_q)$ with injective contraction map, so that $\Phi_f : \mathbb{P}(E^\vee) \rightarrow S_{n-1}$ is defined everywhere. A one-step lower package abstracts the R6 argument: a lower bad locus, the R5 ordered-root curve with its deletions, and the finite set of markers where that construction is unavailable. It consists of:

- (i) a closed bad scheme $B_{n-1} \subset S_{n-1}$ and a finite stratification $S_{n-1} \setminus B_{n-1} = \coprod_\alpha X_\alpha$;
- (ii) for every $\lambda \in \mathbb{P}(E^\vee)(\mathbb{F}_q)$ with $b = \iota_\lambda f \in X_\alpha(\mathbb{F}_q)$, a specified proper geometrically integral curve $Y_{\alpha,b}$, defined over $\mathbb{F}_q$, with genus bound g_α and point bound

$$\#Y_{\alpha,b}(\mathbb{F}_q) \geq q + \kappa_\alpha - 2g_\alpha\sqrt{q};$$

together with a zero-dimensional closed deletion scheme $D_{\alpha,b,\lambda} \subset Y_{\alpha,b}$, of degree at most δ_α , containing the singular, branch, diagonal, marker, and collision points, such that every rational point outside $D_{\alpha,b,\lambda}$ determines an ordered rational split squarefree $g \in W_b$ with $\lambda \nmid g$;

- (iii) a finite parameter scheme $Z_f \subset \mathbb{P}(E^\vee)$, of degree at most d , containing $\Phi_f^{-1}(B_{n-1})$ and every parameter at which the preceding construction is unavailable, including parameters where a fixed factor of the lower system equals a retained marker.

The index α includes any constant-field or Frobenius twist; the assertion in (iii) is the required descent statement. All degrees are geometric degrees on the displayed curve or parameter line.

The construction is summarized in Figure 1. The upper branch is the level-specific component calculation; the lower branch is the general transverse argument.




Fig. 1. Coherent polar induction. The contained branch is a level-specific component theorem; only the transverse branch is controlled by the general point-count argument.

Theorem V.5 (Polar-flag construction). The maps of Definition V.2 commute with field extension and the natural $\text{GL}(E)$ action. They include the infinity chart displayed above. If $g \in W_{\iota_{\lambda_j} \dots \iota_{\lambda_1} f}$, then $\lambda_1 \dots \lambda_j g \in W_f$; the lift is squarefree exactly when g is squarefree and avoids every marker λ_i .

Proof. Contraction is defined by the perfect divided-power/symmetric-power pairing, so it commutes with base change and change of basis. Repeated application of (1) gives the membership assertion. Since the λ_i are distinct on the

configuration space, their product is squarefree; the product with g is squarefree precisely when g is squarefree and none of the λ_i divides it. The coordinate formula at infinity follows by evaluating C_f on the second basis covector. $\square$

Lemma V.6 (Three readings of the collision divisor). Assume W_f has trivial gcd, and let $\lambda \in \mathbb{P}(E^\vee)$. The following are equivalent:

- (i) every member of W_f divisible by λ is divisible by λ^2 ;
- (ii) λ is a fixed factor of $W_{\iota_\lambda f}$;
- (iii) λ lies on the collision divisor of the linear series W_f .

Proof. Equation (1) gives the equality of subspaces

$$\lambda W_{\iota_\lambda f} = W_f \cap \lambda \operatorname{Sym}^{n-2} E^\vee.$$

Thus (i) and (ii) are the same statement after dividing by λ . Since W_f is base-point-free, (i) says exactly that the morphism defined by W_f has zero differential at λ , which is (iii). $\square$

Definition V.7 (One-step contained-component assertion). For a specified bad scheme B_{n-1} , collision divisor, and explicit closed list $\mathcal{C}_n \subset S_n$, the assertion $\operatorname{CC}(n, 1; \mathcal{C}_n)$ says that every f for which the contraction map is noninjective and hence supplies no polar line, $\Phi_f^* B_{n-1}$ is the whole parameter line and hence supplies no transverse marker, or the collision divisor is the whole line and hence every lift repeats its marked root, belongs to $\mathcal{C}_n$. When the list is clear we write $\operatorname{CC}(n, 1)$.

This is a level-specific theorem obligation, not an assumption hidden inside the point-count argument. Section VI states the contained-component calculations used at redundancies six and seven. No assertion $\operatorname{CC}(n, 1)$ is made for an arbitrary degree.

Theorem V.8 (One-step polar escape). Let $f \in S_n(\mathbb{F}_q)$ have a one-step lower package as in Definition V.4. Put

$$\mathcal{H}_\kappa(g, \delta) = 1 + \left\lfloor \left(g + \sqrt{g^2 + \delta - \kappa} \right)^2 \right\rfloor$$

and assume $\delta_\alpha \geq \kappa_\alpha - g_\alpha^2$ for every α . If

$$q \geq \max_\alpha \mathcal{H}_{\kappa_\alpha}(g_\alpha, \delta_\alpha) \quad \text{and} \quad q + 1 > d,$$

then W_f contains a split squarefree member.

Proof. Since $\#\mathbb{P}^1(\mathbb{F}_q) = q + 1 > d$, choose $\lambda \notin Z_f(\mathbb{F}_q)$. On the corresponding lower stratum the recorded point bound gives

$$q + \kappa_\alpha - 2g_\alpha\sqrt{q} > \delta_\alpha,$$

so there is a rational point outside all deletions. It yields a split squarefree lower member avoiding λ . Equation (1) lifts it to a split squarefree member upstairs, so the original system is shallow. $\square$

The theorem isolates the numerical escape and one-step lifting argument. Iterating it requires the contained-component assertion at every intermediate level, not only at the top.

Recursively, the current contraction is chosen outside $\mathcal{P}_j \cup \mathcal{M}_j$; the next CC assertion makes its polar line transverse, so a new marker exists off the finite pullback. No inverse-image stability of the declared loci is assumed.

Remark V.9 (Point-count convention). The correction κ records the exact lower bound available for the chosen model. Every genus-one curve retained in this paper uses the Aubry–Perret bound with $\kappa = 1$; a possible rational singular point is included instead in the deletion degree.

Theorem V.10 (Uniform transverse threshold). Fix a redundancy $r \geq 6$. At every intermediate redundancy $6 \leq j \leq r$, assume the assertion

$$\operatorname{CC}(j - 1, 1; \mathcal{P}_j \cup \mathcal{M}_j)$$

for the recursively pointed lower bad scheme, where $\mathcal{P}_j$ is the persistent rank-two locus and $\mathcal{M}_j$ is the declared nonshallow modular list. Let m be the maximum degree of the stagewise modular pullbacks, taking zero for an empty pullback, and assume they are finite. If

$$m < 3r - 6 + \lfloor 2\sqrt{6r - 17} \rfloor \quad \text{and} \quad q \geq Q_r := 6r - 15 + \lfloor 2\sqrt{6r - 17} \rfloor,$$

every split-free redundancy- r syndrome belongs to $\mathcal{P}_r \cup \mathcal{M}_r$.

Proof. Iterate the pointed contraction to the R5 cubic pencil. Its off-diagonal curve always has $(g, \kappa) = (1, 1)$, and its singular, diagonal, and branch deletions cost thirteen. Each of the $r - 5$ retained markers removes at most six further ordered pairs. Hence

$$\delta_r = 13 + 6(r - 5) = 6r - 17, \quad \mathcal{H}_1(1, \delta_r) = 6r - 15 + \lfloor 2\sqrt{6r - 17} \rfloor.$$

At a stage of redundancy j with $s = r - j$ old markers, the ordinary transverse term costs at most 3, the modular pullback costs at most m , the moving g_{j-2}^{j-4} collision divisor has degree $2j - 6$, and each old marker costs at most three. Thus every marker-selection scheme has degree at most

$$d_j \leq 3 + m + (2j - 6) + 3s = 3r + m - j - 3 \leq 3r + m - 9.$$

The final inequality is sharpest at the redundancy-six marker selection. The displayed hypothesis on m gives $d_j < Q_r$, so the parameter-line inequality $q + 1 > d_j$ is automatic. The exact-linear-gcd graph has only eight base deletions and at most two per marker, so it is weaker as well.

Apply Theorem V.8 successively. If a recursively pointed bad scheme contains the whole polar line, the corresponding assertion CC($j - 1, 1$) places the syndrome in the declared persistent or modular list. Otherwise a marker remains outside its finite scheme. After reaching the R5 pencil, the displayed point bound supplies a split squarefree cubic avoiding every marker, and Theorem V.5 lifts it to the original system. $\square$

For $r = 5, 6, 7$, the integer cutoffs supplied by the same formula are

$$Q_5 = 22, \quad Q_6 = 29, \quad Q_7 = 37;$$

the first prime powers meeting them are 23, 29, 37. Thus the former threshold table consisted of three evaluations of one formula with $Q_r = 6r + O(\sqrt{r})$. At redundancies six and seven the modular pullback degree is $m = 4$, recovering the verified bounds 13 and 16 (the displayed conditions are $m < 20$ and $m < 25$). In general the parameter line remains nonbinding under the modular-degree hypothesis; the other arbitrary-redundancy obstruction is the contained-component chain. Propositions VI.7 and Corollary VI.14 discharge the required chains at redundancies six and seven.

B. Contained-component calculations

The ordinary persistent locus admits a uniform contained-line calculation. This does not classify other lower bad schemes, but it closes the persistent part of every level-specific component gate.

Proposition V.11 (Contained rank-two polar lines). Let $n \geq 5$, let $f \in \mathbb{P}(\Gamma^n E)$, and let

$$A = C_f^{(2)} : \text{Sym}^{n-2} E^\vee \longrightarrow \Gamma^2 E$$

be the three-row catalecticant. Assume the contraction map $E^\vee \rightarrow \Gamma^{n-1} E$ is injective, so the first-polar map is defined on all of $\mathbb{P}(E^\vee)$. This is exactly the complement of the rank-one boundary: after a change of coordinates, a nonzero kernel vector is ∂_X , and $\partial_X f = 0$ forces f to be a scalar multiple of $Y^{[n]}$. Such a point lies on the normal rational curve and is already excluded from the split-free locus. Then

$$\Phi_f(\mathbb{P}(E^\vee)) \subset \{b : \text{rank } C_b^{(2)} \leq 2\} \iff \text{rank } A \leq 2$$

scheme-theoretically and in every characteristic. If $\text{rank } A = 3$, the rank-drop parameters form a subscheme of $\mathbb{P}(E^\vee)$ of degree at most three.

Proof. For a nonzero $\lambda \in E^\vee$, multiplication identifies $\text{Sym}^{n-3} E^\vee$ with the hyperplane

$$H_\lambda = \lambda \text{Sym}^{n-3} E^\vee \subset \text{Sym}^{n-2} E^\vee,$$

and functoriality gives

$$C_{\iota_\lambda f}^{(2)} = A|_{H_\lambda}.$$

The forward implication is immediate. Conversely, suppose $\text{rank } A = 3$ and put $K = \ker A$, so $\dim K = n - 4$. If the lower rank is at most two, rank-nullity on the $(n - 2)$ -dimensional space H_λ gives

$$\dim(K \cap H_\lambda) \geq (n - 2) - 2 = n - 4 = \dim K.$$

Thus $K \subset H_\lambda$. Containment for every geometric $[\lambda]$ would imply

$$K \subset \bigcap_{[\lambda] \in \mathbb{P}(E_k^\vee)} \lambda \text{Sym}^{n-3} E_k^\vee = 0,$$

because a nonzero binary form cannot be divisible by every linear form. This contradicts $\dim K = n - 4 > 0$.

The maximal minors of $C_{\iota_\lambda f}^{(2)}$ are homogeneous cubics in λ . Vanishing at every geometric point is therefore identical vanishing, proving the scheme-theoretic assertion. In the noncontained case at least one cubic is nonzero, and the common rank-drop scheme lies in its degree-three zero divisor. $\square$

Thus Proposition V.11 covers the entire non-rank-boundary locus on which the first-polar line is defined; the excluded noninjective locus is the rank-one normal-rational-curve locus just identified. After rank one and rational split secants are removed as shallow, the upper rank-two locus consists of the persistent tangent and conjugate-secant families. On a persistent fiber the nonsplit torus acts by $z \mapsto z^n$, giving T/T^n modulo inversion and Frobenius, while the tangent unipotent cocycle is $z \mapsto z + nu$.

Modular loci are equally explicit. If $\mathcal{N}$ is a nucleus in the lower divided-power module, the complete space whose first-polar line lies in $\mathcal{N}$ is

$$\mathcal{M}_n(\mathcal{N}) = \mathbb{P} \ker \left(\Gamma^n E \longrightarrow \text{Hom}(E^\vee, \Gamma^{n-1} E / \tilde{\mathcal{N}}) \right).$$

Lucas' theorem computes the nucleus coordinates, and this displayed kernel computes every coherent lift without an ambient syndrome census.

Remark V.12 (Scope of the contained calculations). The modular contraction kernel is an exact finite linear-algebra construction. Proposition V.11 supplies the ordinary persistent component uniformly, but other lower bad schemes and collision divisors still require level-specific calculations. Neither the proposition nor Theorem V.8 classifies those additional components.

VI. Redundancies six and seven: complete and gated results

A. Redundancy six

For

$$f = (a_0 : \cdots : a_5) \in \mathbb{P}(\Gamma^5 E), \quad W_f = \ker \begin{pmatrix} a_0 & a_1 & a_2 & a_3 & a_4 \\ a_1 & a_2 & a_3 & a_4 & a_5 \end{pmatrix},$$

the system W_f is a net of quartics off the quintic normal rational curve. It is split-free exactly when it has no completely split squarefree quartic.

Proposition VI.1 (Persistent net and orbit law). If $\gcd(W_f) = Q$ has degree two, then

$$W_f = Q \text{Sym}^2 E^\vee.$$

The split-free persistent locus has size $q(q+1)^2/2$. On the tangent fiber the unipotent coordinate is $x \mapsto x + 5u$; on the irreducible quadratic fiber the norm-one torus acts through $z \mapsto z^5$, with inversion and coefficient Frobenius acting on T/T^5 .

Proof. For fixed Q , the syndromes annihilating $Q \text{Sym}^3 E^\vee$ form a projective line. If Q is irreducible, all $q+1$ points remain; if $Q = \lambda^2$, one endpoint has dependent Hankel rows, leaving q points. Hence

$$q(q+1) + \frac{q(q-1)}{2}(q+1) = \frac{q(q+1)^2}{2}.$$

Normalize a square to U^2 . The nondegenerate fiber is $[e_4 + xe_5]$, and direct degree-five substitution gives $x \mapsto x + 5u$. For an irreducible Q , its stabilizer identifies the $q+1$ -point fiber with $T = \{z \in \mathbb{F}_{q^2}^\times : z^{q+1} = 1\}$ and acts by fifth powers; the normalizer inverts z , while coefficient Frobenius multiplies T/T^5 by the characteristic. This proves the stated orbit law and the characteristic-five tangent splitting. $\square$

Lemma VI.2 (Exact linear gcd is shallow). Let $q \geq 7$. If $\gcd(W_f)$ has degree one, then W_f contains a split squarefree quartic.

Proof. Write $W_f = LV$, where L is rational and $V \subset \text{Sym}^3 E^\vee$ has vector dimension three. It suffices to find in V a product of three distinct rational factors, none equal to L . Move the root of L to infinity and let ℓ span $V^\perp$. On the cubic with distinct finite roots x, y, z , write

$$\ell((T - xU)(T - yU)(T - zU)) = c_0 + c_1(x + y + z) + c_2(xy + xz + yz) + c_3xyz =: P(x, y, z).$$

Suppose that this value is nonzero for every distinct triple. For fixed distinct x, y , write $P = A(x, y) + zB(x, y)$, where

$$B(x, y) = c_1 + c_2(x + y) + c_3xy.$$

If $B(x, y) \neq 0$, the unique zero in z must be x or y . Consequently

$$B(x, y)P(x, y, x)P(x, y, y) = 0$$

for every $x \neq y$; it also vanishes when $B(x, y) = 0$. This polynomial has degree at most four in each variable. Since $q - 1 > 4$, first varying $y \neq x$ and then x shows that it is the zero polynomial. The polynomial ring is a domain. Neither $P(x, y, x)$ nor $P(x, y, y)$ vanishes identically unless $\ell = 0$, as comparison of the coefficients of $1, x, y, x^2, xy, x^2y$ shows. Hence $B = 0$, so $c_1 = c_2 = c_3 = 0$.

Thus V is the hyperplane of cubics vanishing at infinity, and every member of V is divisible by L . This would make $\gcd(W_f)$ divisible by L^2 , contrary to its exact degree. Therefore a required cubic exists, and multiplying it by L gives a split squarefree member of W_f . $\square$

Proposition VI.3 (First-polar line and secant degree). For $r \in \mathbb{F}_q$ put

$$b(r) = (a_1 - ra_0, \dots, a_5 - ra_4), \quad b(\infty) = (a_0, \dots, a_4).$$

Then

$$(T - rU)g \in W_f \iff g \in W_{b(r)}.$$

If W_f has trivial gcd, its polar line meets the lower secant cubic in degree at most three.

Proof. The first equivalence follows by expanding the two Hankel equations of $(T - rU)g$. Let

$$C_f = \begin{pmatrix} a_0 & a_1 & a_2 & a_3 \\ a_1 & a_2 & a_3 & a_4 \\ a_2 & a_3 & a_4 & a_5 \end{pmatrix}.$$

The lower catalecticant equals

$$C_f \begin{pmatrix} -r & 0 & 0 \\ 1 & -r & 0 \\ 0 & 1 & -r \\ 0 & 0 & 1 \end{pmatrix}.$$

Cauchy–Binet gives the secant determinant

$$D_f(r) = -\Delta_{012}r^3 + \Delta_{013}r^2 - \Delta_{023}r + \Delta_{123}.$$

It vanishes identically exactly when $\text{rank } C_f \leq 2$, equivalently when W_f has quadratic gcd. In the trivial-gcd case it therefore has at most three zeros. $\square$

Proposition VI.4 (Cyclic and collision degrees). For a trivial-gcd net:

- (i) in characteristic different from two and three, the lower tame cyclic carrier is a binomially rescaled Veronese surface of degree four and contains no polar line;
- (ii) the pointed ramification condition has degree at most six and does not vanish identically;
- (iii) an S_3 lower pencil with the marked root deleted has total deletion degree at most nineteen.

Proof. In syndrome coordinates the tame cyclic surface is

$$\{\text{diag}(1, 4, 6, 4, 1)^{-1}[Q^2] : [Q] \in \mathbb{P}(\text{Sym}^2 E^\vee)\}.$$

When 2 and 3 are invertible this is the Veronese surface, of degree four, and contains no line. Characteristic three has a different wild closure, treated in Proposition VI.5. For a basis F_0, F_1, F_2 of the $g_4^2 = W_f$, the collision divisor is the rank-one condition on the value and derivative rows. Equivalently, all three binary Wronskians

$$M_{ij} = F_i dF_j - F_j dF_i \quad (0 \leq i < j \leq 2)$$

vanish at the marker. Each M_{ij} is a binary form of degree six: on an affine chart the apparent degree-seven leading terms cancel, and homogenization includes the point at infinity.

At least one M_{ij} is nonzero. Otherwise the morphism $\mathbb{P}^1 \rightarrow \mathbb{P}^2$ defined by the base-point-free net W_f would have zero differential and hence factor through absolute Frobenius. In characteristic three this is impossible because its pullback line bundle has degree four, not divisible by three. In characteristic two the factorization would identify W_f , after a change of coordinates, with the full pure-square space $\langle T^4, T^2U^2, U^4 \rangle$. Its annihilator is $\langle e_1^\vee, e_3^\vee \rangle$. For the two consecutive Hankel rows this would require simultaneously

$$a_0 = a_2 = a_4 = 0, \quad a_1 = a_3 = a_5 = 0,$$

contrary to $f \neq 0$. Thus the collision scheme is contained in the zero divisor of one nonzero M_{ij} , and has degree at most six in every characteristic.

Finally the R5 off-diagonal curve has deletion degree thirteen, including its possible rational singular point. The unique cubic through a specified marker contributes at most six ordered pairs, giving $(g, \delta, \kappa) = (1, 19, 1)$ and

$$q + 1 - 2\sqrt{q} > 19.$$

Its first prime-power solution is 29. $\square$

Proposition VI.5 (Modular contained components). In characteristic three the lower wild carrier is the degree-four cone

$$\text{Join}(e_2, C_4),$$

and no polar line of a trivial-gcd net is contained in it. In characteristic two the cyclic carrier is the plane

$$\Pi_{\text{cyc}} = \mathbb{P}\langle e_1, e_2, e_3 \rangle,$$

whose unique contained polar-line component is

$$\mathcal{N}_3 = \mathbb{P}\langle e_2, e_3 \rangle.$$

Proof. Every line on the wild cone is a ruling. By equivariance it suffices to compare the two consecutive rows on one ruling; in the frame $\langle e_2, e_4 \rangle$ the conditions

$$(a_0, \dots, a_4), (a_1, \dots, a_5) \in \langle e_2, e_4 \rangle$$

force the overlapping coordinates successively to vanish. Undoing the frame gives $f = \mu\nu(t)$, including $f = \mu e_5$ at infinity. These are fixed-factor/rank-boundary syndromes, so no trivial-gcd polar line is contained in the wild cone.

In characteristic two, direct substitution gives

$$g \cdot e_2 = [0 : \delta\gamma : \alpha\delta + \beta\gamma : \alpha\beta : 0],$$

whose closure is Π_{cyc} . The consecutive-row overlap is now literal: the first row lies in the plane exactly when $a_0 = a_4 = 0$, and the second exactly when $a_1 = a_5 = 0$. Hence a polar line lies in the plane exactly when

$$a_0 = a_1 = a_4 = a_5 = 0,$$

namely when $f \in \mathbb{P}\langle e_2, e_3 \rangle$. This proves both existence and uniqueness of the contained binary component. $\square$

Proposition VI.6 (Binary nucleus arithmetic). The scheme $\mathcal{N}_3(\mathbb{F}_q)$ is one projective-semilinear orbit of $q + 1$ points. It is split-free over $q = 2^m$ exactly when m is odd.

Proof. For the representative e_3 ,

$$W_{e_3} = \langle 1, t, t^4 \rangle.$$

The root differences of $A + Bt + Ct^4$ lie in the kernel of $u \mapsto u^4 + (B/C)u$. Four distinct rational roots exist exactly when all roots of $u^3 = B/C$ are rational, equivalently when $3 \mid q - 1$, which for $q = 2^m$ means m is even. The stabilizer of e_3 is the Borel of order $q(q - 1)$, so the orbit has $q + 1$ points; its lower-unipotent transforms and endpoint are precisely $\mathcal{N}_3(\mathbb{F}_q)$. $\square$

Proposition VI.7 (Exhaustive R6 contained-component classification). For a redundancy-six syndrome, the alternatives in Definition V.7 are exhausted as follows:

- (i) a noninjective contraction map is the rank-one normal-rational-curve boundary and is shallow;
- (ii) a polar line contained in the lower secant cubic comes from the upper rank-two catalecticant locus; after its shallow rational secants are removed, this is the persistent tangent and conjugate-secant locus;
- (iii) no trivial-gcd polar line is contained in the tame cyclic surface or the characteristic-three wild cone;
- (iv) in characteristic two, containment in the cyclic plane occurs exactly for $f \in \mathcal{N}_3 = \mathbb{P}\langle e_2, e_3 \rangle$;
- (v) on the remaining trivial-gcd locus the collision divisor is not identically zero.

Consequently the complete nonshallow contained list is the persistent locus together with $\mathcal{N}_3$ in characteristic two. There is no unlisted contained component.

Proof. The noninjective and secant-cubic alternatives are Proposition V.11 and Proposition VI.3. The tame cyclic and collision assertions are Proposition VI.4; the wild and binary specializations are Proposition VI.5. These are all the clauses of Definition V.7: contraction indeterminacy, containment in each component of the lower bad scheme, and identical collision. Hence the list is exhaustive. $\square$

Proposition VI.8 (R6 one-step exhaustion). Assume $q \geq 29$. Then every split-free redundancy-six syndrome lies in the persistent quadratic-gcd locus or, when $q = 2^m$ with odd m , in $\mathcal{N}_3(\mathbb{F}_q)$.

Proof. The complete contained branch is Proposition VI.7. The common-factor trichotomy is exhaustive. A quadratic gcd gives the persistent tangent and conjugate-secant loci, together with the shallow rational secant; an exact linear gcd is shallow by Lemma VI.2. It remains to treat a trivial-gcd net.

Its first-polar line is not contained in the lower secant cubic, by Proposition VI.3. The pullback of that cubic has degree at most three. The tame cyclic surface has degree four and contains no polar line. In characteristic three its

wild replacement has the same degree and no contained trivial-gcd polar line; in characteristic two the replacement is a plane, whose only contained component is $\mathcal{N}_3$, by Proposition VI.5. Finally, Proposition VI.4 gives a nonzero collision divisor of degree at most six. Outside the contained binary component, the excluded parameter scheme therefore has degree

$$d \leq 3 + 4 + 6 = 13.$$

For every remaining marker, the lower cubic pencil has exact linear gcd or geometric monodromy S_3 . In the first case the graph argument of Proposition IV.3 is pointed as follows: besides its eight original deletions, remove the at most two graph points whose split quadratic contains the new marker. Since $q + 1 > 10$, a split squarefree cubic avoiding that marker remains. Here the fixed gcd λ_0 differs from the marker r . Indeed, $\lambda_0 = r$ says that every cubic in $W_{b(r)}$ vanishes at r , so every lifted quartic has a double root there; by Lemma V.6, this is exactly a point of the degree-six collision divisor already placed in Z_f , so no additional deletion budget is needed. In the S_3 case, Lemma IV.6 supplies the geometrically integral off-diagonal curve. Its singular, diagonal, and branch deletions have degree at most thirteen, and deleting the unique cubic through the marker removes at most six ordered pairs. Thus

$$(g, \delta, \kappa) = (1, 19, 1), \quad q + 1 - 2\sqrt{q} > 19$$

for every $q \geq 29$. Explicitly,

$$\mathcal{H}_1(1, 19) = 1 + \left\lfloor (1 + \sqrt{19})^2 \right\rfloor = 29, \quad 30 - 2\sqrt{29} > 19.$$

These data form the one-step package of Definition V.4, with $d = 13$, so Theorem V.8 produces a split squarefree quartic.

This is one field-order threshold, not three characteristic-dependent estimates: the first characteristic-three and characteristic-two prime powers above it happen to be 81 and 32, respectively. On $\mathcal{N}_3$, Proposition VI.6 gives the asserted odd-degree split-free orbit and proves shallowness in even extension degree. $\square$

The covering-radius gate is complete: Seroussi–Roth gives $\rho(\text{PRS}(q - 5)) = 5$ for $q \geq 9$, while Certificate R6 checks it directly at $q = 7, 8$.

Theorem VI.9 (Redundancy six). For every prime power $q \geq 7$, the deep syndromes of $\text{PRS}(q - 5)$ are the persistent locus of size $q(q + 1)^2/2$, together with:

- (i) respectively 18, 11, 4, 2, 1 exceptional $\text{PGL}_2(q)$ orbits at $q = 7, 8, 9, 11, 13$;
- (ii) one characteristic-two nucleus orbit of size $q + 1$ when $q = 2^m$ and odd $m \geq 5$.

There are no other exceptional orbits. The complete persistent orbit law is T/T^5 modulo inversion and Frobenius, with a tangent fixed/nonzero split exactly in characteristic five. At $q = 7, 8, 9, 11, 13$, the exceptional direction totals are

$$5152, \quad 4713, \quad 1800, \quad 792, \quad 546,$$

and the corresponding total numbers of deep syndrome directions are

$$5376, \quad 5037, \quad 2250, \quad 1584, \quad 1820.$$

Proof. Propositions VI.7 and VI.8 give the complete conceptual exhaustion for every prime power $q \geq 29$. Certificate R6 exhausts

$$q \in \{7, 8, 9, 11, 13, 16, 17, 19, 23, 25, 27\},$$

finds exceptional projective-orbit counts 18, 11, 4, 2, 1 only at the five displayed fields, and no further orbit at the remaining fields. Together with the separate radius gate, this covers every prime power $q \geq 7$. $\square$

The listed sources are disjoint. The persistent locus has quadratic gcd, every finite exceptional row has trivial gcd, and the separate nucleus term occurs only for binary fields beyond the finite table.

The five small tables are exhaustive computations. For an exceptional net W_f , let $I_f(r)$ be the number of irreducible cubic members of the quotient pencil $W_{b(r)}$, including the infinity chart, and define its shared-root pair count by

$$E_f = \sum_{r \in \mathbb{P}^1(\mathbb{F}_q)} \binom{I_f(r)}{2}.$$

In odd characteristic let

$$F_f(T, U) = \sum_{i=0}^5 a_i T^{5-i} U^i$$

be the binary quintic attached to $f = (a_0 : \cdots : a_5)$, and let $\tau_5(f)$ be its rational factorization type; for example, $1^3 + 2$ denotes a triple rational factor and an irreducible quadratic factor. Their semilinear orbit counts are 18, 5, 2, 2, 1; E_f ,

together in odd characteristic with $\tau_5(f)$, separates the classes up to precisely Frobenius fusion. Table VI prints the canonical representatives, stabilizers, orbit sizes, and fusion lengths.

TABLE VI: Exceptional redundancy-six classes. Each row is one projective-semilinear class. The columns give a canonical syndrome representative, the size and stabilizer order of each constituent $\mathrm{PGL}_2(q)$ -orbit, the number ϕ of such orbits fused by coefficient Frobenius, the shared-root pair count E , and the binary-quintic factor type τ_5 in odd characteristic.

q	representative	$ \mathcal{O} $	$ G_f $	ϕ	E	τ_5
7	$[0 : 0 : 0 : 1 : 0 : 1]$	168	2	1	7	$1^3 + 2$
7	$[0 : 0 : 1 : 0 : 3 : 1]$	336	1	1	19	$1^2 + 1 + 2$
7	$[0 : 0 : 1 : 0 : 3 : 2]$	336	1	1	15	$1^2 + 3$
7	$[0 : 1 : 0 : 0 : 0 : 1]$	168	2	1	10	$1 + 2 + 2$
7	$[0 : 1 : 0 : 0 : 2 : 0]$	112	3	1	33	$1 + 1 + 3$
7	$[0 : 1 : 0 : 1 : 1 : 2]$	336	1	1	9	$1 + 4$
7	$[0 : 1 : 0 : 1 : 1 : 5]$	336	1	1	9	$1 + 1 + 3$
7	$[0 : 1 : 0 : 1 : 3 : 0]$	336	1	1	21	$1 + 1 + 1 + 2$
7	$[0 : 1 : 0 : 1 : 3 : 3]$	336	1	1	7	$1 + 1 + 3$
7	$[0 : 1 : 0 : 1 : 3 : 6]$	336	1	1	14	$1 + 1 + 3$
7	$[0 : 1 : 0 : 3 : 0 : 5]$	168	2	1	11	$1 + 4$
7	$[0 : 1 : 0 : 3 : 0 : 6]$	168	2	1	20	$1 + 4$
7	$[0 : 1 : 0 : 3 : 2 : 1]$	336	1	1	13	$1 + 1 + 3$
7	$[1 : 0 : 0 : 0 : 1 : 2]$	336	1	1	7	$1 + 4$
7	$[1 : 0 : 0 : 0 : 1 : 3]$	336	1	1	23	5
7	$[1 : 0 : 0 : 3 : 1 : 3]$	336	1	1	13	$1 + 4$
7	$[1 : 0 : 1 : 0 : 5 : 2]$	336	1	1	8	$1 + 4$
7	$[1 : 0 : 1 : 0 : 6 : 2]$	336	1	1	8	5
8	$[0 : 0 : 0 : 1 : 0 : 0]$	9	56	1	0	—
8	$[0 : 1 : 1 : 0 : 2 : 1]$	504	1	3	20	—
8	$[0 : 1 : 1 : 0 : 2 : 5]$	504	1	3	19	—
8	$[1 : 0 : 0 : 1 : 3 : 6]$	504	1	3	14	—
8	$[1 : 0 : 1 : 3 : 5 : 2]$	168	3	1	54	—
9	$[0 : 1 : 0 : 0 : 0 : 4]$	180	4	2	28	$1 + 4$
9	$[0 : 1 : 0 : 1 : 4 : 2]$	720	1	2	29	$1 + 4$
11	$[0 : 0 : 0 : 1 : 0 : 0]$	132	10	1	60	$1^3 + 1^2$
11	$[0 : 1 : 0 : 2 : 0 : 10]$	660	2	1	30	$1 + 4$
13	$[0 : 1 : 0 : 0 : 0 : 2]$	546	4	1	60	$1 + 4$

For $q = 8$, integers are binary polynomial-basis codes with $\alpha^3 = \alpha + 1$; for $q = 9$, they are ternary polynomial-basis codes with $\alpha^2 = -1$. Thus $\sum \phi = (18, 11, 4, 2, 1)$ and $\sum \phi|\mathcal{O}| = (5152, 4713, 1800, 792, 546)$, in the field order 7, 8, 9, 11, 13.

B. Redundancy seven

For

$$f = (a_0 : \dots : a_6) \in \mathbb{P}(\Gamma^6 E), \quad W_f = \ker \begin{pmatrix} a_0 & a_1 & a_2 & a_3 & a_4 & a_5 \\ a_1 & a_2 & a_3 & a_4 & a_5 & a_6 \end{pmatrix},$$

the system W_f is a four-dimensional web of quintics. For an affine marker r the contractions are

$$b(r) = (a_1 - ra_0, \dots, a_6 - ra_5).$$

Coefficient collection gives

$$(t - r)g \in W_f \iff g \in W_{b(r)};$$

the lift is squarefree exactly when $g(r) \neq 0$.

The contracted net $W_{b(r)}$ has quadratic gcd exactly when its three-row catalecticant has rank at most two. Unless the first-polar line is contained in that rank-two locus, these parameters form the degree-at-most-three intersection of Proposition VI.3; the contained case is Corollary VI.14.

Proposition VI.10 (Complete two-marker lower package). Let W_h be a trivial-gcd quartic net and let x be a prescribed forbidden root. Assume in characteristic two that $h \notin \mathcal{N}_3$. For every prime power $q \geq 37$, the net contains a split squarefree quartic avoiding x .

Proof. By equivariance choose coordinates in which x is finite; all parameter equations below are then homogenized in s , so the infinity chart is included. Choose the second contraction marker s . The unavailable values of s comprise: $s = x$, of degree one; at most three intersections with the lower secant cubic, by Proposition VI.3; at most four intersections with the cyclic surface or its characteristic-three wild replacement; and at most six points of the collision divisor, by Propositions VI.4 and VI.5. In characteristic two the cyclic replacement is a plane, so its pullback has degree at most one unless the polar line is contained in it, the excluded case $h \in \mathcal{N}_3$, by Proposition VI.5. We must also exclude the parameters for which the lower pencil has fixed factor x . They are cut out by the 3-minors obtained by appending the fixed evaluation row $(1, x, x^2, x^3)$ to the two Hankel rows of $b(s)$, and hence have degree at most two in s . These minors are not all zero. If they vanished identically, then

$$W_h \cap \{g : g(s) = 0\} \subseteq \{g : g(x) = 0\}$$

for every geometric s . Given $g \in W_h$, choose a geometric root s of g ; then $g(x) = 0$. Thus x would divide every member of W_h , contradicting the trivial-gcd hypothesis. Hence the complete second-step parameter scheme has degree at most

$$1 + 3 + 4 + 6 + 2 = 16 < q + 1.$$

Choose s outside it.

The resulting cubic pencil is neither persistent nor cyclic. If it has exact linear gcd λ_0 , then $\lambda_0 \notin \{x, s\}$: equality with x is excluded by the fixed-factor minors above, while equality with s is the pointed collision excluded by Proposition VI.4. The graph proof of Proposition IV.3 has eight original deletions; avoiding the two additional markers costs at most two graph points each. Since $q + 1 > 12$, a split squarefree cubic avoiding x and s remains.

Otherwise the bottom cover has geometric monodromy S_3 . Its off-diagonal curve is geometrically integral of arithmetic genus one. The singular, diagonal, and branch deletions cost at most thirteen, and the fibers through x and s cost at most six ordered pairs each. Hence

$$(g, \delta, \kappa) = (1, 25, 1), \quad q + 1 - 2\sqrt{q} > 25,$$

whose first prime-power solution is 37. In either case the cubic lifts through s to a split squarefree quartic in W_h avoiding x . $\square$

Proposition VI.11 (Exact-gcd-one avoidance). Suppose the lower net is $\ell_x U$, where $U = \ker \lambda \subset \text{Sym}^3 E^\vee$ has dimension three. If $x = r$, the parameter belongs to the collision divisor. If $x \neq r$ and $q \geq 16$, then U contains a split cubic avoiding both x and r .

Proof. If $x = r$, every lifted quintic has the repeated factor $(t - r)^2$. Lemma V.6 identifies this with the degree-eight collision divisor of Proposition VI.12. This is the unavailable branch, so assume $x \neq r$. Put $S = \mathbb{P}^1(\mathbb{F}_q) \setminus \{x, r\}$. For ordered distinct $u, v \in S$, the map $w \mapsto \lambda(\ell_u \ell_v \ell_w)$ is linear. If it vanishes identically, choose any remaining w ; otherwise it has one rational zero. The equations $w = x, r, u, v$ have respective bidegrees $(1, 1), (1, 1), (2, 1), (1, 2)$ in (u, v) . None is identically zero: the first two would add a common factor, and the last two are excluded because double-root cubics span $\text{Sym}^3 E^\vee$. A nonzero bidegree- (a, b) divisor on $\mathbb{P}^1 \times \mathbb{P}^1$ has at most $(a + b)(q + 1)$ rational points. Thus the bad pairs number at most $10(q + 1)$, whereas $(q - 1)(q - 2) > 10(q + 1)$ for $q \geq 16$. $\square$

Proposition VI.12 (Collision divisor). After fixed factors are removed, the collision divisor of the moving g_5^3 has degree at most eight.

Proof. It is the ramification divisor of a separable g_5^3 , whose degree is $(3 + 1)(5 - 3) = 8$. If the series were inseparable, it would factor through Frobenius and lie in the p -th-power subspace of quintics; that subspace has dimension at most three, smaller than the four-dimensional Hankel series. $\square$

The characteristic-two nucleus line at redundancy six has exactly one coherent lift, the central point e_3 , with

$$W_{e_3} = \langle 1, t, t^4, t^5 \rangle.$$

It is deep for odd m ; for even m , the homogenization of $t^4 + t$ has the five roots $\mathbb{P}^1(\mathbb{F}_4)$.

Proposition VI.13 (Central binary lift). The complete inverse image of the R6 nucleus line under consecutive sextic contraction is $[e_3]$. It is split-free over $\mathbb{F}_{2^m}$ exactly when m is odd.

Proof. Putting both consecutive rows in $\mathbb{P}\langle e_2, e_3 \rangle$ gives $a_0 = a_1 = a_2 = a_4 = a_5 = a_6 = 0$. Thus only e_3 remains, with $W_{e_3} = \langle 1, t, t^4, t^5 \rangle$. For odd m , contraction at a root of a hypothetical split quintic would give a split quartic on the R6 nucleus line, contradicting Proposition VI.6. For even m , $\mathbb{F}_4 \subset \mathbb{F}_q$ and the homogenization of $t^4 + t$ has its four affine $\mathbb{F}_4$ roots and the simple root at infinity. $\square$

Corollary VI.14 (Exhaustive R7 contained-component classification). The assertion CC(6, 1) is exhaustive in every characteristic. A noninjective contraction is the shallow rank-one boundary. Containment in the lower persistent locus gives upper catalecticant rank at most two; after the shallow rational secants are removed, these are exactly the persistent tangent and conjugate-secant families. The complete lift of the only additional lower component, the characteristic-two nucleus line, is the central point e_3 . The remaining collision divisor is nonzero. Thus the nonshallow contained list is precisely the persistent locus and, in characteristic two, e_3 ; there is no further contained component.

Proof. Apply Proposition V.11 with $n = 6$. The three rational types of the quadratic recurrence are split squarefree, repeated rational, and irreducible. They give respectively the shallow rational secant, tangent, and conjugate-secant loci. The remaining lower binary nucleus line has the separate complete lift computed in Proposition VI.13. Proposition VI.12 shows that the collision divisor is nonzero. These are respectively the contraction-indeterminacy, lower-component-containment, and identical-collision alternatives of Definition V.7, so the list is exhaustive. $\square$

Theorem VI.15 (Redundancy-seven classification). Certificate R7 unconditionally classifies every prime power $q \leq 32$ in its displayed domain. For every prime power $q \geq 13$, the split-free syndromes at redundancy seven are the persistent locus of size $q(q+1)^2/2$, together with the central nucleus point when $q = 2^m$ and m is odd. At $q = 7, 8, 9, 11$ the additional split-free orbit-size profiles, after removing the persistent locus and also the central nucleus singleton at $q = 8$, are:

$$\begin{array}{l|l} 7 & 56, 84^5, 112^2, 168^{45}, 336^{141} \\ 8 & 63, 72, 84^3, 168^4, 252^{24}, 504^{86} \\ 9 & 180^3, 240^6, 360^{18}, 720^{27} \\ 11 & 264^2, 440, 660^2. \end{array}$$

This is the complete all-field split-free classification. Since Seroussi–Roth gives $\rho(\text{PRS}(q-6)) = 6$ for every prime power $q \geq 11$, it is also the complete deep-hole classification in that range. At $q = 7, 8, 9$ the table is not promoted to a code deep-hole classification without a separate covering-radius calculation. The persistent orbit law is T/T^6 modulo inversion and Frobenius, with tangent splitting in characteristics two and three.

Proof. For fixed irreducible quadratic Q , the persistent fiber has $q+1$ points, while for $Q = \lambda^2$ it has q non-rank-one points. There are $q(q-1)/2$ irreducible quadratics and $q+1$ rational squares. Hence the persistent count is

$$\frac{q(q-1)}{2}(q+1) + (q+1)q = \frac{q(q+1)^2}{2}.$$

For $q \geq 37$, choose a first marker outside the at most three catalecticant intersections, eight collision parameters, and, in characteristic two, one intersection with the nucleus line. If the contracted quartic net has trivial gcd, Proposition VI.10 supplies the complete second one-step package: it excludes the cyclic/wild carrier and fixed-gcd/marker coincidences before applying either the pointed linear-gcd graph or the S_3 cover. If the contracted net has exact linear gcd, Proposition VI.11 applies; its excluded $x = r$ case is already among the eight collision parameters. Corollary VI.14 leaves exactly the persistent locus and Proposition VI.13. Certificate R7 covers

$$q \in \{7, 8, 9, 11, 13, 16, 17, 19, 23, 25, 27, 29, 31, 32\}$$

by an independent five-secant replay. The radius theorem promotes the split-free list only from $q \geq 11$. $\square$

For the four finite rows Certificate R7 contains the canonical representative, stabilizer, persistent/central flags, and Frobenius link of every orbit. The displayed size profile is only a compact summary; repeated sizes are distinguished by those canonical records. There are respectively 194, 119, 54, 5 exceptional representatives at $q = 7, 8, 9, 11$; the complete lists are printed in supplement/CLASSIFICATION-RECORDS.md, while the matching JSON is the machine-readable

record. The following gives one printed canonical representative for every orbit-size block in the displayed profile; repeated orbits within a block are distinguished in the full list:

q	$ \mathcal{O} $	representative
7	56	$[0 : 0 : 0 : 0 : 1 : 0 : 0]$
	84	$[0 : 1 : 0 : 0 : 0 : 1 : 0]$
	112	$[0 : 0 : 1 : 0 : 0 : 2 : 0]$
	168	$[0 : 0 : 0 : 0 : 1 : 0 : 1]$
	336	$[0 : 0 : 0 : 1 : 0 : 1 : 1]$
8	63	$[0 : 0 : 0 : 1 : 0 : 0 : 1]$
	72	$[0 : 0 : 0 : 0 : 1 : 0 : 0]$
	84	$[0 : 1 : 1 : 1 : 1 : 3 : 6]$
	168	$[1 : 0 : 1 : 1 : 3 : 7 : 7]$
	252	$[0 : 0 : 1 : 0 : 1 : 0 : 0]$
	504	$[0 : 0 : 1 : 0 : 1 : 0 : 1]$
9	180	$[0 : 1 : 0 : 0 : 0 : 4 : 0]$
	240	$[0 : 0 : 1 : 0 : 1 : 1 : 1]$
	360	$[0 : 1 : 0 : 1 : 3 : 3 : 2]$
	720	$[0 : 0 : 1 : 1 : 0 : 2 : 2]$
11	264	$[0 : 1 : 0 : 0 : 0 : 0 : 2]$
	440	$[1 : 0 : 1 : 1 : 4 : 3 : 3]$
	660	$[1 : 0 : 1 : 0 : 1 : 4 : 8]$.

For $q = 8, 9$, the coordinates use the same polynomial-basis encoding specified below Table VI.

The $q = 19$ lower net

$$W = \langle 1, t^3, t^4 \rangle$$

shows why the parameter must remain marked. Its split squarefree members are exactly

$$U(T^3 - uU^3), \quad u \in (\mathbb{F}_{19}^\times)^3 = \{1, 7, 8, 11, 12, 18\}.$$

All six contain the marked point at infinity. The obstruction is pointed rather than intrinsic to W : at the marker 0, for example, $U(T^3 - U^3)$ is an admissible member. No coherent sextic polar line lies in the infinity-pointed orbit. Thus the R7 proof escapes over the marker parameter; classifying unpointed bad lower fibers would retain a false exception.

TABLE VII

Field-range closure ledger. “Certificate” means a complete finite classification with canonical representatives, not merely an orbit count.

Result	prime powers	closure mechanism
R5	7, 8, 9, 11, 13, 16, 17, 19 $q \geq 23$	direct census and independent replay cubic-cover geometry, Lemma IV.6
R6	7, 8, 9, 11, 13, 16 17, 19, 23, 25, 27 $q \geq 29$	direct syndrome/radius census finite structural bridge certificate marked-cover point bound and contained components
R7	7, 8, 9, 11 13, 16, 17, 19, 23, 25, 27, 29, 31, 32 $q \geq 37$	direct split-free census coherent-polar finite bridge certificate marked-cover point bound and contained components

The R7 rows classify split-free syndromes. Conversion to code deep-hole directions uses the separate covering-radius gate and is made only for $q \geq 11$. The public certificate schema and exact replay boundaries are given in the supplement; the table does not promote orbit-size profiles to identifiers.

There is no prime power in any of the open numerical intervals $(19, 23)$, $(27, 29)$, or $(32, 37)$. Thus adjacent finite and geometric rows in Table VII leave no unlisted field order.

VII. Verification and formal boundary

No geometric classification theorem in this paper is claimed to be formalized in Lean. The formal layer checks coordinate identities and conditional synthesis implications with the component geometry, rational-point existence, group actions, and certificate semantics kept as explicit inputs.

The conceptual theorems above are accompanied by deterministic evidence bundles. Table VIII records which claims use exhaustive computation. Each public label resolves, through the supplement schema, to its generator, certificate, replay where stated, and SHA-256 manifest. Table I gives the corresponding claim-level mathematical dependencies.

TABLE VIII: Public verification map. Exact filenames, commands, hashes, domains, and stop conditions belong to the electronic supplement.

Public label	Verified domain	Replay and trust boundary
Certificate R5	nineteen finite fields	independent pointwise and orbitwise replay; canonical representatives and exhaustion records
Certificate R6	direct scan through $q = 16$; finite bridge below 29	independent syndrome replay and separate radius gate
Certificate R6-NF	small exceptional normal forms	deterministic normal-form reconstruction; not a second implementation
Certificate R7	$q = 7, 8, 9, 11$ and finite bridge below 37	independent five-secant/orbit checks; split-free and radius flags remain separate

Table VI is a deterministic projection of the hash-pinned public R6 and R6-NF records. The top-level verifier regenerates it and checks byte-for-byte equality with the included source.

The public labels are defined by `supplement/CERTIFICATE-SCHEMA.md`. The paper-local bundle, literal replay commands, toolchain locks, hashes, byte counts, and working directories are recorded in `supplement/REPRODUCING.md`, `supplement/EVIDENCE-MANIFEST.json`, and `supplement/RELEASE-MANIFEST.md`. The single entry point `python3 supplement/verify.py` checks the bundle manifest, classification records and their expected hashes, and the local PDF and `supplement-artifact` rows printed in the release manifest; its `-replay` option runs the paper-local Python replays.

The paper-facing Lean gate checks coordinate algebra and conditional synthesis, not the geometric classifications or rational-point arguments. The exact declarations, hypotheses, axiom audit, and reproduction command are listed in `supplement/LEAN-STATEMENTS.md`.

Data, code, and disclosure

The electronic supplement accompanying this manuscript contains the source, canonical classification representatives and stabilizers, certificate schemas, generators, independent replays where available, expected hashes, toolchain locks, and a top-level verification command. Section VII assigns each computational or formal claim its exact trust route.

VIII. Scope and open problems

The first levels and the contained/transverse fork point to one stable picture. The bottom-cover point count is uniform and linear: Theorem V.10 gives its exact threshold, subject to an explicit nonbinding bound on the modular pullback degree. The possible quadratic scale below is therefore allowance for the contained and modular component problem, not for rational points on the bottom cover.

Conjecture VIII.1 (Stable polar classification). There is an absolute constant C such that, at redundancy r , every split-free projective syndrome over $\mathbb{F}_q$ with $q \geq Cr^2$ lies either in the persistent locus or in a nonshallow component of the level- r Lucas contraction kernel. Equivalently, transverse arithmetic exceptions disappear uniformly, and the contained locus has no further stable component. In every characteristic for which that Lucas kernel has no nonshallow component, the persistent tangent and conjugate-secant families are the complete high-field list.

The Lucas clause is forced by the binary nucleus families at redundancies six and seven, which persist over infinitely many fields. The conjecture predicts that rank-two persistence and digit-controlled contraction kernels exhaust the stable locus, with the coherent polar flag eliminating everything transverse.

The present results establish the first levels of this picture but do not constitute an arbitrary-redundancy classification.

- 1) Redundancy five is complete for every $q \geq 7$.
- 2) Redundancy six is a complete all-field deep-hole classification. Redundancy seven has a complete all-field split-free classification; its $q = 7, 8, 9$ covering-radius gate remains separate.
- 3) At higher redundancy the remaining input is the chain of level-specific contained-component assertions. The bottom monodromy and point count are the uniform R5 calculation.

Thus Theorem V.10 is conditional rather than an arbitrary-redundancy classification: higher-level persistent and modular contained loci remain open here. In particular, the paper does not settle the general Reed–Solomon deep-hole conjecture.

The structural conclusion is nevertheless stable. Beyond redundancy four, deepness is no longer controlled by one invariant or by isolated bad fibers. The persistent geometry is a coherent rank-two polar flag; modular geometry appears through characteristic-specific contraction kernels; and the remaining arithmetic is the Frobenius action on a normalized ordered-root cover. Redundancy five is the first exceptional-cover case, while redundancies six and seven show how that cover is propagated or contained.

Appendix

Acknowledgment

OpenAI Codex was used under the author's direction throughout the project for proof exploration, symbolic and finite computation, code generation, and drafting assistance. Drafting assistance affected prose throughout the manuscript and electronic supplement. The author checked the resulting arguments and computations and assumes responsibility for every claim in the manuscript. The system is not an author.

References

- [1] K. Kaipa, “Deep holes and MDS extensions of Reed–Solomon codes,” *IEEE Transactions on Information Theory*, vol. 63, no. 8, pp. 4940–4948, Aug. 2017.
- [2] J. Zhang, D. Wan, and K. Kaipa, “Deep holes of projective Reed–Solomon codes,” *IEEE Transactions on Information Theory*, vol. 66, no. 4, pp. 2392–2401, 2020.
- [3] Y. Aubry and M. Perret, “Coverings of singular curves over finite fields,” *Manuscripta Mathematica*, vol. 88, pp. 467–478, 1995.
- [4] G. Seroussi and R. M. Roth, “On MDS extensions of generalized Reed–Solomon codes,” *IEEE Transactions on Information Theory*, vol. 32, no. 3, pp. 349–354, 1986.
- [5] S. Ball and M. Lavrauw, “Arcs in finite projective spaces,” *EMS Surveys in Mathematical Sciences*, vol. 6, no. 1–2, pp. 133–172, 2019.
- [6] A. Dür, “The automorphism groups of Reed–Solomon codes,” *Journal of Combinatorial Theory, Series A*, vol. 44, no. 1, pp. 69–82, 1987.
- [7] J. Zhang and D. Wan, “Explicit deep holes of Reed–Solomon codes,” <https://arxiv.org/abs/1711.02292>, 2017, version 1, 7 November 2017.
- [8] X. Xu, S. Hong, and Y. Xu, “On deep holes of generalized projective Reed–Solomon codes,” <https://arxiv.org/abs/1705.07823>, 2017, version 1, 22 May 2017.
- [9] X. Xu, “On deep holes of projective Reed–Solomon codes over finite fields with even characteristic,” *Wuhan University Journal of Natural Sciences*, vol. 28, no. 1, pp. 15–19, 2023.
- [10] Y. Wu, C. Ding, and T. Chen, “Extended codes and deep holes of MDS codes,” <https://arxiv.org/abs/2312.05534>, 2023, version 1, 9 December 2023.
- [11] A. Iarrobino and V. Kanev, *Power Sums, Gorenstein Algebras, and Determinantal Loci*, ser. *Lecture Notes in Mathematics*. Berlin, Heidelberg: Springer, 1999, vol. 1721.
- [12] G. Comas and M. Seiguer, “On the rank of a binary form,” *Foundations of Computational Mathematics*, vol. 11, pp. 65–78, 2011.
- [13] E. Cesaratto, G. Matera, and M. Pérez, “The distribution of factorization patterns on linear families of polynomials over a finite field,” *Combinatorica*, vol. 37, no. 5, pp. 805–836, 2017.
- [14] J. Gmainer and H. Havlicek, “Nuclei of normal rational curves,” *Journal of Geometry*, vol. 69, pp. 117–130, 2000.
- [15] T. Wang, “Splitting of polynomial families via galois theory,” <https://arxiv.org/abs/2606.12810>, 2026, version 1, 11 June 2026.